


Accuracy vs. latency — Llama-3.1-8B-Instruct, $B = 2048$, 32K input



Both axes measured in one harness, all 11 InfiniteBench tasks. Latency: 32K input / 512 generated, batch 1, RTX PRO 6000 Blackwell.
Accuracy: 50 seeded-random records per task, greedy. FullKV is dense (no page selection), so it has no latency point.