



KV-compression memory — measured breakdown (left) and exact scaling (right); RTX A6000 48GB

What each method costs
Qwen-2.5-7B, 32K in, $\beta=2048$



peak alloc separates the arms
by only 4.8% (25603-26838 MB)
-- why it is not the headline

How it scales
ours costs 1.91x min/max per page



Peak GPU memory is not the headline: the harness preallocates fixed 6 GiB gpu / 20 GiB cpu arenas regardless of method, so it separates the arms by <5%.
One page-summary row per method (min/max, compact min/max, rank-8 basis) counted as PAGES ACTUALLY USED. Our records are int4/int8 (LOCKS' r8i4 scheme): 978 B per page per kv head. RaaS is O(1) in context; the others are O(context).