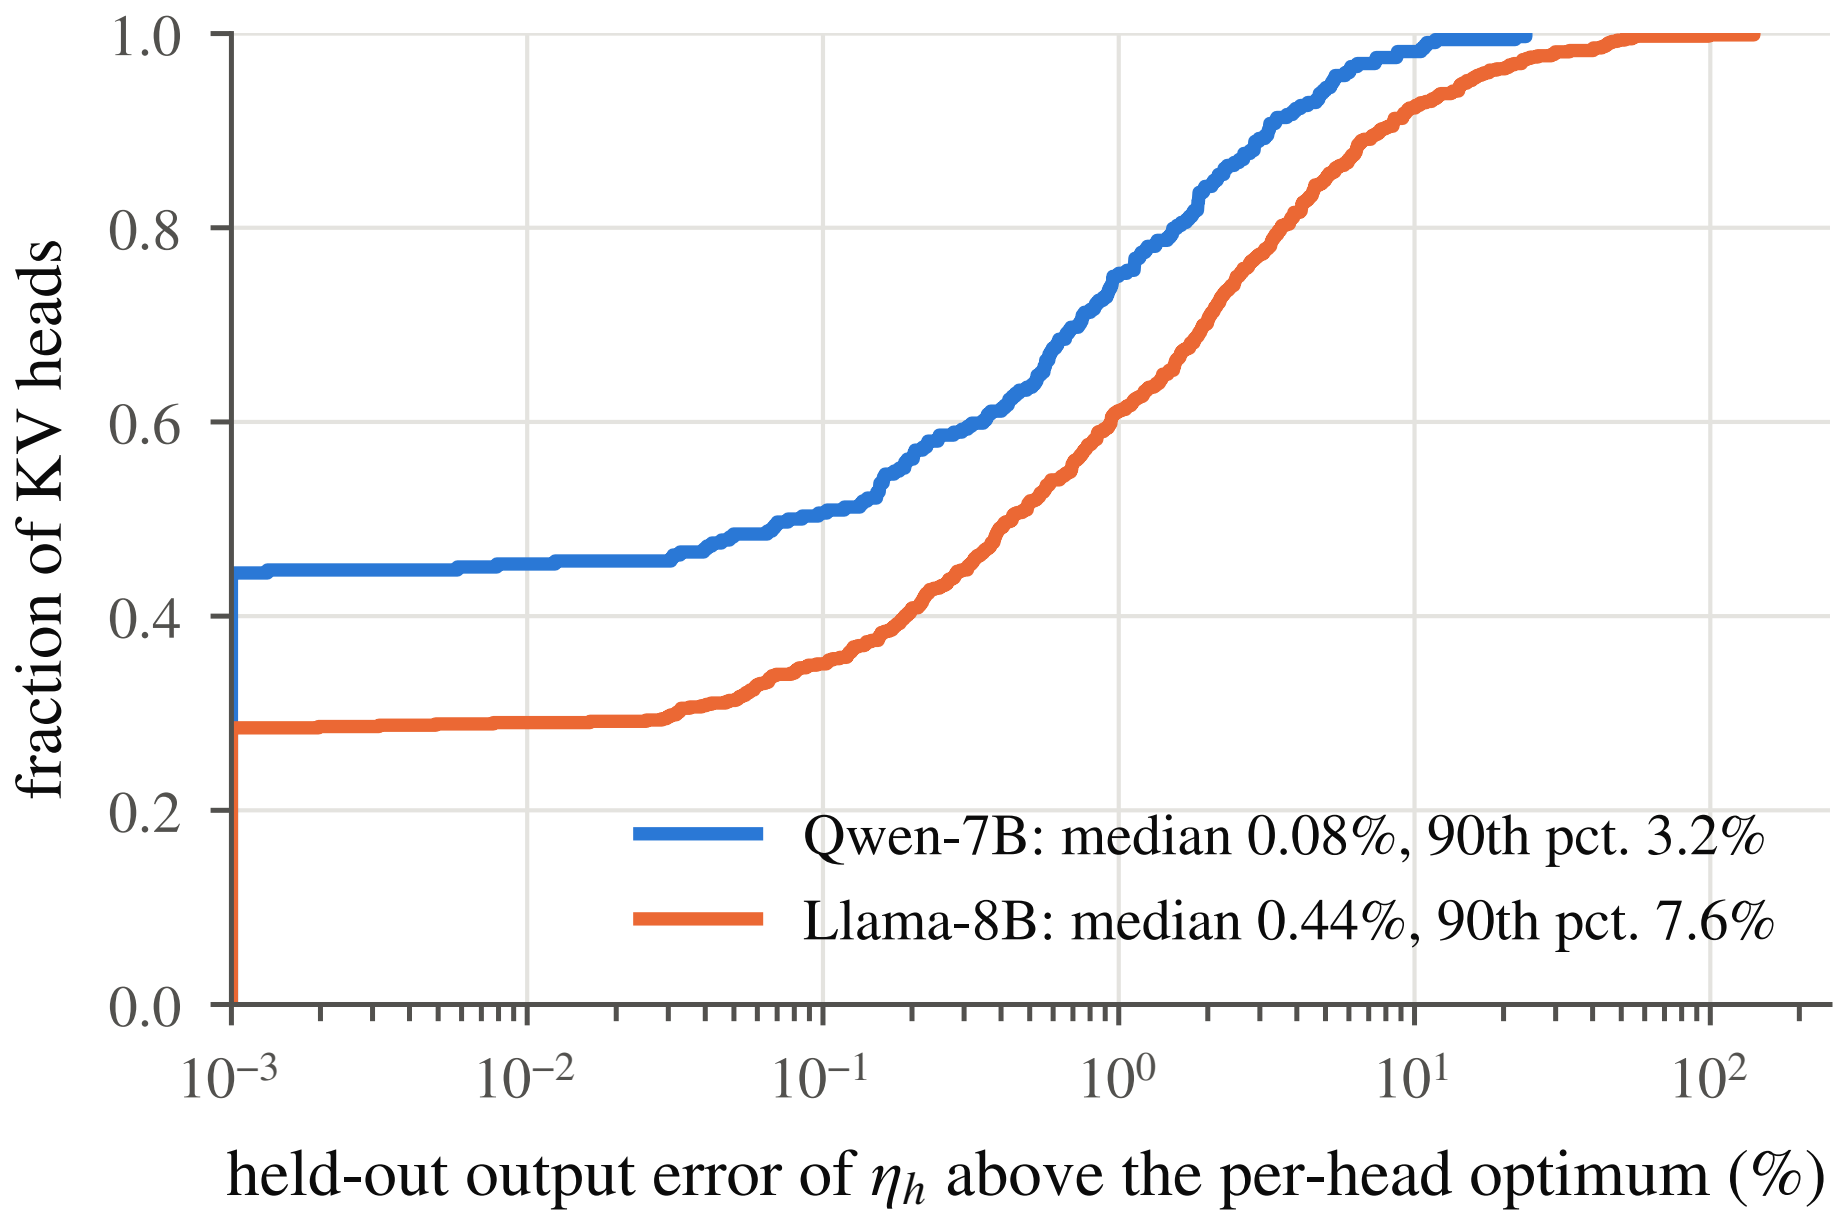
1.0
0.8
0.6
0.4
0.2
0.0
fraction of KV heads
Qwen-7B: median 0.08%, 90th pct. 3.2%
Llama-8B: median 0.44%, 90th pct. 7.6%
10^{-3}
10^{-2}
10^{-1}
10^{0}
10^{1}
10^{2}
held-out output error of η_h above the per-head optimum (%)